

```
self.out_tensor = out_tensor
return out_tensor
```



tf.keras and tf.estimator

TensorFlow has now integrated the popular Keras object-oriented frontend into the core TensorFlow library. Keras includes a `Layer` class definition that closely matches the `Layer` objects you’ve just learned about in this section. In fact, the `Layer` objects here were adapted from the DeepChem library, which in turn adapted them from an earlier version of Keras.

It’s worth noting, though, that `tf.keras` has not yet become the standard higher-level interface to TensorFlow. The `tf.estimator` module provides an alternative (albeit less rich) high-level interface to raw TensorFlow.

Regardless of which frontend eventually becomes standard, we think that understanding the fundamental design principles for building your own frontend is instructive and worthwhile. You might need to build a new system for your organization that requires an alternative design, so a solid grasp of design principles will serve you well.

Defining a Graph of Layers

We mentioned briefly in the previous section that a deep architecture could be visualized as a directed graph of `Layer` objects. In this section, we transform this intuition into the `TensorGraph` object. These objects are responsible for constructing the underlying TensorFlow computation graph.

A `TensorGraph` object is responsible for maintaining a `tf.Graph`, a `tf.Session`, and a list of layers (`self.layers`) internally (Example 8-9). The directed graph is represented implicitly, by the `in_layers` belonging to each `Layer` object. `TensorGraph` also contains utilities for saving this `tf.Graph` instance to disk and consequently assigns itself a directory (using `tempfile.mkdtemp()` if none is specified) to store checkpoints of the weights associated with its underlying TensorFlow graph.

Example 8-9. The `TensorGraph` contains a graph of layers; `TensorGraph` objects can be thought of as the “model” object holding the deep architecture you want to train

```
class TensorGraph(object):

    def __init__(self,
                 batch_size=100,
                 random_seed=None,
                 graph=None,
```

```

        learning_rate=0.001,
        model_dir=None,
        **kwargs):
    """
    Parameters
    -----
    batch_size: int
        default batch size for training and evaluating
    graph: tensorflow.Graph
        the Graph in which to create Tensorflow objects. If None, a new Graph
        is created.
    learning_rate: float or LearningRateSchedule
        the learning rate to use for optimization
    kwargs
    """

    # Layer Management
    self.layers = dict()
    self.features = list()
    self.labels = list()
    self.outputs = list()
    self.task_weights = list()
    self.loss = None
    self.built = False
    self.optimizer = None
    self.learning_rate = learning_rate

    # Singular place to hold Tensor objects which don't serialize
    # See TensorGraph._get_tf() for more details on lazy construction
    self.tensor_objects = {
        "Graph": graph,
        #"train_op": None,
    }
    self.global_step = 0
    self.batch_size = batch_size
    self.random_seed = random_seed
    if model_dir is not None:
        if not os.path.exists(model_dir):
            os.makedirs(model_dir)
    else:
        model_dir = tempfile.mkdtemp()
        self.model_dir_is_temp = True
    self.model_dir = model_dir
    self.save_file = "%s/%s" % (self.model_dir, "model")
    self.model_class = None

```

The private method `_add_layer` does bookkeeping work to add a new Layer object to the `TensorGraph` (Example 8-10).

Example 8-10. The `_add_layer` method adds a new `Layer` object

```
def _add_layer(self, layer):
    if layer.name is None:
        layer.name = "%s_%s" % (layer.__class__.__name__, len(self.layers) + 1)
    if layer.name in self.layers:
        return
    if isinstance(layer, Input):
        self.features.append(layer)
    self.layers[layer.name] = layer
    for in_layer in layer.in_layers:
        self._add_layer(in_layer)
```

The layers in a `TensorGraph` must form a directed acyclic graph (there can be no loops in the graph). As a result, we can topologically sort these layers. Intuitively, a topological sort “orders” the layers in the graph so that each `Layer` object’s `in_layers` precede it in the ordered list. This topological sort is necessary to make sure all input layers to a given layer are added to the graph before the layer itself ([Example 8-11](#)).

Example 8-11. The `topsort` method orders the layers in the `TensorGraph`

```
def topsort(self):

    def add_layers_to_list(layer, sorted_layers):
        if layer in sorted_layers:
            return
        for in_layer in layer.in_layers:
            add_layers_to_list(in_layer, sorted_layers)
        sorted_layers.append(layer)

    sorted_layers = []
    for l in self.features + self.labels + self.task_weights + self.outputs:
        add_layers_to_list(l, sorted_layers)
    add_layers_to_list(self.loss, sorted_layers)
    return sorted_layers
```

The `build()` method takes the responsibility of populating the `tf.Graph` instance by calling `layer.create_tensor` for each layer in topological order ([Example 8-12](#)).

Example 8-12. The `build` method populates the underlying `TensorFlow` graph

```
def build(self):
    if self.built:
        return
    with self._get_tf("Graph").as_default():
        self._training_placeholder = tf.placeholder(dtype=tf.float32, shape=())
        if self.random_seed is not None:
            tf.set_random_seed(self.random_seed)
        for layer in self.topsort():
```

```

        with tf.name_scope(layer.name):
            layer.create_tensor(training=self._training_placeholder)
self.session = tf.Session()

self.built = True

```

The method `set_loss()` adds a loss for training to the graph. `add_output()` specifies that the layer in question might be fetched from the graph. `set_optimizer()` specifies the optimizer used for training (Example 8-13).

Example 8-13. These methods add necessary losses, outputs, and optimizers to the computation graph

```

def set_loss(self, layer):
    self._add_layer(layer)
    self.loss = layer

def add_output(self, layer):
    self._add_layer(layer)
    self.outputs.append(layer)

def set_optimizer(self, optimizer):
    """Set the optimizer to use for fitting."""
    self.optimizer = optimizer

```

The method `get_layer_variables()` is used to fetch the learnable `tf.Variable` objects created by a layer. The private method `_get_tf` is used to fetch the `tf.Graph` and optimizer instances underpinning the `TensorGraph`. `get_global_step` is a convenience method for fetching the current step in the training process (starting from 0 at construction). See Example 8-14.

Example 8-14. Fetch the learnable variables associated with each layer

```

def get_layer_variables(self, layer):
    """Get the list of trainable variables in a layer of the graph."""
    if not self.built:
        self.build()
    with self._get_tf("Graph").as_default():
        if layer.variable_scope == "":
            return []
        return tf.get_collection(
            tf.GraphKeys.TRAINABLE_VARIABLES, scope=layer.variable_scope)

def get_global_step(self):
    return self._get_tf("GlobalStep")

def _get_tf(self, obj):
    """Fetches underlying TensorFlow primitives.

```

```

Parameters
-----
obj: str
    If "Graph", returns tf.Graph instance. If "Optimizer", returns the
    optimizer. If "train_op", returns the train operation. If "GlobalStep" returns
    the global step.
Returns
-----
TensorFlow Object
"""

if obj in self.tensor_objects and self.tensor_objects[obj] is not None:
    return self.tensor_objects[obj]
if obj == "Graph":
    self.tensor_objects["Graph"] = tf.Graph()
elif obj == "Optimizer":
    self.tensor_objects["Optimizer"] = tf.train.AdamOptimizer(
        learning_rate=self.learning_rate,
        beta1=0.9,
        beta2=0.999,
        epsilon=1e-7)
elif obj == "GlobalStep":
    with self._get_tf("Graph").as_default():
        self.tensor_objects["GlobalStep"] = tf.Variable(0, trainable=False)
return self._get_tf(obj)

```

Finally, the `restore()` method restores a saved `TensorGraph` from disk (Example 8-15). (As you will see later, the `TensorGraph` is saved automatically during training.)

Example 8-15. Restore a trained model from disk

```

def restore(self):
    """Reload the values of all variables from the most recent checkpoint file."""
    if not self.built:
        self.build()
    last_checkpoint = tf.train.latest_checkpoint(self.model_dir)
    if last_checkpoint is None:
        raise ValueError("No checkpoint found")
    with self._get_tf("Graph").as_default():
        saver = tf.train.Saver()
        saver.restore(self.session, last_checkpoint)

```

The A3C Algorithm

In this section you will learn how to implement A3C, the asynchronous reinforcement learning algorithm you saw earlier in the chapter. A3C is a significantly more complex training algorithm than those you have seen previously. The algorithm requires running gradient descent in multiple threads, interspersed with game rollout