
Table of Contents

Preface.....	ix
1. Introduction to Deep Learning.....	1
Machine Learning Eats Computer Science	1
Deep Learning Primitives	3
Fully Connected Layer	3
Convolutional Layer	4
Recurrent Neural Network Layers	4
Long Short-Term Memory Cells	5
Deep Learning Architectures	6
LeNet	6
AlexNet	6
ResNet	7
Neural Captioning Model	8
Google Neural Machine Translation	9
One-Shot Models	10
AlphaGo	12
Generative Adversarial Networks	13
Neural Turing Machines	14
Deep Learning Frameworks	15
Limitations of TensorFlow	16
Review	17
2. Introduction to TensorFlow Primitives.....	19
Introducing Tensors	19
Scalars, Vectors, and Matrices	20
Matrix Mathematics	24
Tensors	25

Tensors in Physics	27
Mathematical Asides	28
Basic Computations in TensorFlow	29
Installing TensorFlow and Getting Started	29
Initializing Constant Tensors	30
Sampling Random Tensors	31
Tensor Addition and Scaling	32
Matrix Operations	33
Tensor Types	35
Tensor Shape Manipulations	35
Introduction to Broadcasting	37
Imperative and Declarative Programming	37
TensorFlow Graphs	39
TensorFlow Sessions	39
TensorFlow Variables	40
Review	42
3. Linear and Logistic Regression with TensorFlow.....	43
Mathematical Review	43
Functions and Differentiability	44
Loss Functions	45
Gradient Descent	50
Automatic Differentiation Systems	53
Learning with TensorFlow	55
Creating Toy Datasets	55
New TensorFlow Concepts	60
Training Linear and Logistic Models in TensorFlow	64
Linear Regression in TensorFlow	64
Logistic Regression in TensorFlow	73
Review	79
4. Fully Connected Deep Networks.....	81
What Is a Fully Connected Deep Network?	81
“Neurons” in Fully Connected Networks	83
Learning Fully Connected Networks with Backpropagation	85
Universal Convergence Theorem	87
Why Deep Networks?	88
Training Fully Connected Neural Networks	89
Learnable Representations	89
Activations	89
Fully Connected Networks Memorize	90
Regularization	90

Training Fully Connected Networks	94
Implementation in TensorFlow	94
Installing DeepChem	94
Tox21 Dataset	95
Accepting Minibatches of Placeholders	96
Implementing a Hidden Layer	96
Adding Dropout to a Hidden Layer	97
Implementing Minibatching	98
Evaluating Model Accuracy	98
Using TensorBoard to Track Model Convergence	99
Review	101
5. Hyperparameter Optimization.....	103
Model Evaluation and Hyperparameter Optimization	104
Metrics, Metrics, Metrics	105
Binary Classification Metrics	106
Multiclass Classification Metrics	108
Regression Metrics	110
Hyperparameter Optimization Algorithms	110
Setting Up a Baseline	111
Graduate Student Descent	113
Grid Search	114
Random Hyperparameter Search	115
Challenge for the Reader	116
Review	117
6. Convolutional Neural Networks.....	119
Introduction to Convolutional Architectures	120
Local Receptive Fields	120
Convolutional Kernels	122
Pooling Layers	125
Constructing Convolutional Networks	125
Dilated Convolutions	126
Applications of Convolutional Networks	127
Object Detection and Localization	127
Image Segmentation	128
Graph Convolutions	129
Generating Images with Variational Autoencoders	131
Training a Convolutional Network in TensorFlow	134
The MNIST Dataset	134
Loading MNIST	135
TensorFlow Convolutional Primitives	138

The Convolutional Architecture	140
Evaluating Trained Models	144
Challenge for the Reader	146
Review	146
7. Recurrent Neural Networks.....	149
Overview of Recurrent Architectures	150
Recurrent Cells	152
Long Short-Term Memory (LSTM)	152
Gated Recurrent Units (GRU)	154
Applications of Recurrent Models	154
Sampling from Recurrent Networks	154
Seq2seq Models	155
Neural Turing Machines	157
Working with Recurrent Neural Networks in Practice	159
Processing the Penn Treebank Corpus	159
Code for Preprocessing	160
Loading Data into TensorFlow	162
The Basic Recurrent Architecture	164
Challenge for the Reader	166
Review	166
8. Reinforcement Learning.....	169
Markov Decision Processes	173
Reinforcement Learning Algorithms	175
Q-Learning	176
Policy Learning	177
Asynchronous Training	179
Limits of Reinforcement Learning	179
Playing Tic-Tac-Toe	181
Object Orientation	181
Abstract Environment	182
Tic-Tac-Toe Environment	182
The Layer Abstraction	185
Defining a Graph of Layers	188
The A3C Algorithm	192
The A3C Loss Function	196
Defining Workers	198
Training the Policy	201
Challenge for the Reader	203
Review	203

9. Training Large Deep Networks.....	205
Custom Hardware for Deep Networks	205
CPU Training	206
GPU Training	207
Tensor Processing Units	209
Field Programmable Gate Arrays	211
Neuromorphic Chips	211
Distributed Deep Network Training	212
Data Parallelism	213
Model Parallelism	214
Data Parallel Training with Multiple GPUs on Cifar10	215
Downloading and Loading the DATA	216
Deep Dive on the Architecture	218
Training on Multiple GPUs	220
Challenge for the Reader	223
Review	223
10. The Future of Deep Learning.....	225
Deep Learning Outside the Tech Industry	226
Deep Learning in the Pharmaceutical Industry	226
Deep Learning in Law	227
Deep Learning for Robotics	227
Deep Learning in Agriculture	228
Using Deep Learning Ethically	228
Is Artificial General Intelligence Imminent?	230
Where to Go from Here?	231
Index.....	233