

```

Out[19]: subject      Bob      Guido      Sue
         type      HR      Temp      HR      Temp      HR      Temp
         year visit
2013 1      31.0  38.7  32.0  36.7  35.0  37.2
        2      44.0  37.7  50.0  35.0  29.0  36.7
2014 1      30.0  37.4  39.0  37.8  61.0  36.9
        2      47.0  37.8  48.0  37.3  51.0  36.5

```

Here we see where the multi-indexing for both rows and columns can come in *very* handy. This is fundamentally four-dimensional data, where the dimensions are the subject, the measurement type, the year, and the visit number. With this in place we can, for example, index the top-level column by the person's name and get a full Data Frame containing just that person's information:

```

In[20]: health_data['Guido']

Out[20]: type      HR      Temp
         year visit
2013 1      32.0  36.7
        2      50.0  35.0
2014 1      39.0  37.8
        2      48.0  37.3

```

For complicated records containing multiple labeled measurements across multiple times for many subjects (people, countries, cities, etc.), use of hierarchical rows and columns can be extremely convenient!

Indexing and Slicing a MultiIndex

Indexing and slicing on a MultiIndex is designed to be intuitive, and it helps if you think about the indices as added dimensions. We'll first look at indexing multiply indexed Series, and then multiply indexed DataFrames.

Multiply indexed Series

Consider the multiply indexed Series of state populations we saw earlier:

```

In[21]: pop

Out[21]: state      year
         California 2000    33871648
                2010    37253956
         New York   2000    18976457
                2010    19378102
         Texas      2000    20851820
                2010    25145561
dtype: int64

```

We can access single elements by indexing with multiple terms:

```

In[22]: pop['California', 2000]

Out[22]: 33871648

```

The MultiIndex also supports *partial indexing*, or indexing just one of the levels in the index. The result is another Series, with the lower-level indices maintained:

```
In[23]: pop['California']  
Out[23]: year  
         2000    33871648  
         2010    37253956  
dtype: int64
```

Partial slicing is available as well, as long as the MultiIndex is sorted (see discussion in “Sorted and unsorted indices” on page 137):

```
In[24]: pop.loc['California':'New York']  
Out[24]: state      year  
         California  2000    33871648  
                    2010    37253956  
         New York   2000    18976457  
                    2010    19378102  
dtype: int64
```

With sorted indices, we can perform partial indexing on lower levels by passing an empty slice in the first index:

```
In[25]: pop[:, 2000]  
Out[25]: state  
         California    33871648  
         New York     18976457  
         Texas        20851820  
dtype: int64
```

Other types of indexing and selection (discussed in “Data Indexing and Selection” on page 107) work as well; for example, selection based on Boolean masks:

```
In[26]: pop[pop > 22000000]  
Out[26]: state      year  
         California  2000    33871648  
                    2010    37253956  
         Texas      2010    25145561  
dtype: int64
```

Selection based on fancy indexing also works:

```
In[27]: pop[['California', 'Texas']]  
Out[27]: state      year  
         California  2000    33871648  
                    2010    37253956  
         Texas      2000    20851820  
                    2010    25145561  
dtype: int64
```

Multiply indexed DataFrames

A multiply indexed DataFrame behaves in a similar manner. Consider our toy medical DataFrame from before:

```
In[28]: health_data

Out[28]:
```

		Bob		Guido		Sue	
	type	HR	Temp	HR	Temp	HR	Temp
year visit							
2013	1	31.0	38.7	32.0	36.7	35.0	37.2
	2	44.0	37.7	50.0	35.0	29.0	36.7
2014	1	30.0	37.4	39.0	37.8	61.0	36.9
	2	47.0	37.8	48.0	37.3	51.0	36.5

Remember that columns are primary in a DataFrame, and the syntax used for multiply indexed Series applies to the columns. For example, we can recover Guido’s heart rate data with a simple operation:

```
In[29]: health_data['Guido', 'HR']

Out[29]:
```

year	visit	
2013	1	32.0
	2	50.0
2014	1	39.0
	2	48.0

Name: (Guido, HR), dtype: float64

Also, as with the single-index case, we can use the `loc`, `iloc`, and `ix` indexers introduced in “Data Indexing and Selection” on page 107. For example:

```
In[30]: health_data.iloc[:, :2]

Out[30]:
```

		Bob	
	type	HR	Temp
year visit			
2013	1	31.0	38.7
	2	44.0	37.7

These indexers provide an array-like view of the underlying two-dimensional data, but each individual index in `loc` or `iloc` can be passed a tuple of multiple indices. For example:

```
In[31]: health_data.loc[:, ('Bob', 'HR')]

Out[31]:
```

year	visit	
2013	1	31.0
	2	44.0
2014	1	30.0
	2	47.0

Name: (Bob, HR), dtype: float64

Working with slices within these index tuples is not especially convenient; trying to create a slice within a tuple will lead to a syntax error:

```
In[32]: health_data.loc[:, 1], (:, 'HR')]
```

File "<ipython-input-32-8e3cc151e316>", line 1
 health_data.loc[:, 1], (:, 'HR')]
 ^
SyntaxError: invalid syntax

You could get around this by building the desired slice explicitly using Python's built-in `slice()` function, but a better way in this context is to use an `IndexSlice` object, which Pandas provides for precisely this situation. For example:

```
In[33]: idx = pd.IndexSlice
        health_data.loc[idx[:, 1], idx[:, 'HR']]
```

```
Out[33]: subject      Bob Guido   Sue
         type      HR      HR      HR
         year visit
2013 1      31.0  32.0  35.0
2014 1      30.0  39.0  61.0
```

There are so many ways to interact with data in multiply indexed `Series` and `Data Frames`, and as with many tools in this book the best way to become familiar with them is to try them out!

Rearranging Multi-Indices

One of the keys to working with multiply indexed data is knowing how to effectively transform the data. There are a number of operations that will preserve all the information in the dataset, but rearrange it for the purposes of various computations. We saw a brief example of this in the `stack()` and `unstack()` methods, but there are many more ways to finely control the rearrangement of data between hierarchical indices and columns, and we'll explore them here.

Sorted and unsorted indices

Earlier, we briefly mentioned a caveat, but we should emphasize it more here. *Many of the MultiIndex slicing operations will fail if the index is not sorted.* Let's take a look at this here.

We'll start by creating some simple multiply indexed data where the indices are *not lexicographically sorted*:

```
In[34]: index = pd.MultiIndex.from_product(['a', 'c', 'b'], [1, 2])
        data = pd.Series(np.random.rand(6), index=index)
        data.index.names = ['char', 'int']
        data
```

```
Out[34]: char  int
         a      1      0.003001
         a      2      0.164974
         c      1      0.741650
```